

Retrieval-Augmented Generation (RAG)

RAG is a technique used in Artificial Intelligence to improve the quality of responses generated by Large Language Models. Instead of relying only on trained knowledge, RAG retrieves relevant information from external documents before generating answers. **Steps in RAG:**

1. Document Loading
2. Chunking
3. Embeddings
4. Vector Database
5. Retrieval
6. Response Generation

RAG is widely used in PDF chat applications, research assistants, educational tools, and company knowledge bases.